

DD2476: Search Engines and Information Retrieval Systems

Johan Boye*

KTH

Lecture 6

* Many slides inspired by Manning, Raghavan and Schütze

Improving recall in search

- Relevance feedback (assignment 3.1)
 - “Give me more of this (and less of that)”
- Wildcard queries (3.3, 3.4)
 - “*colo*rful*” → “*colorful*”, “*colourful*”
- Spelling correction (3.5, 3.6)
 - “*see you on the wki*” → “*see you on the wiki*”
- Query expansion
 - adding synonyms, etc. to the query
 - word vectors, multi-lingual retrieval

Relevance feedback

Relevance feedback: Basic idea

- The user issues a (short, simple) query.
- The search engine returns a set of documents.
- User marks some docs as relevant, some as nonrelevant.
- Search engine computes a new query that better (hopefully) represents the information need.
- Search engine runs new query and returns new results.
- New results have (hopefully) better recall.


Key concept for relevance feedback: Centroid

- The centroid is the center of mass of a set of points.
- Recall that we represent documents as points in a high-dimensional space.
- Thus: we can compute centroids of documents.
- Definition:

$$\vec{\mu}(D) = \frac{1}{|D|} \sum_{d \in D} \vec{v}(d)$$

- where D is a set of documents, and $\vec{v}(d) = \vec{d}$ is the vector we use to represent document d .

Key concept for relevance feedback: Centroid



Rocchio algorithm

- The Rocchio algorithm implements relevance feedback in the vector space model.
- Rocchio chooses the query \vec{q}_{opt} that maximizes

$$\vec{q}_{opt} = \arg \max_{\vec{q}} [\text{sim}(\vec{q}, \mu(D_r)) - \text{sim}(\vec{q}, \mu(D_{nr}))]$$

- D_r : set of relevant docs; D_{nr} : set of nonrelevant docs
- Intent: \vec{q}_{opt} is the vector that separates relevant and nonrelevant docs maximally.
- We can rewrite this as:

$$\vec{q}_{opt} = \mu(D_r) + [\mu(D_r) - \mu(D_{nr})]$$


Rocchio algorithm

- The optimal query vector is:

$$\begin{aligned}\vec{q}_{opt} &= \mu(D_r) + [\mu(D_r) - \mu(D_{nr})] \\ &= \frac{1}{|D_r|} \sum_{\vec{d}_j \in D_r} \vec{d}_j + \left[\frac{1}{|D_r|} \sum_{\vec{d}_j \in D_r} \vec{d}_j - \frac{1}{|D_{nr}|} \sum_{\vec{d}_j \in D_{nr}} \vec{d}_j \right]\end{aligned}$$


- We move the centroid of the relevant documents by the difference between the two centroids.

Example: Rocchio algorithm




circles: relevant documents, Xs: nonrelevant documents

Rocchio illustrated




$\vec{\mu}_R$: centroid of relevant documents

Rocchio illustrated




$\vec{\mu}_R$ does not separate relevant / nonrelevant.

Rocchio illustrated




$\vec{\mu}_R$ centroid of nonrelevant documents.

Rocchio illustrated




$\vec{\mu}_R - \vec{\mu}_{NR}$: difference vector

Rocchio illustrated




Add difference vector to $\vec{\mu}_R$...

Rocchio illustrated



... to get \vec{q}_{opt}

Rocchio illustrated



\vec{q}_{opt} separates relevant / nonrelevant perfectly.

Quiz

Four documents:

1. cat ✓

2. cat dog ✓

3. cat horse horse ✗

4. horse ✗

I find **1,2 relevant**, and **3,4 non-relevant**. What is the optimal query according to Rocchio?

Quiz

1	✓	cat	(1, 0, 0)
2	✓	cat dog	(1, 1, 0)
3	✗	cat horse horse	(1, 0, 2)
4	✗	horse	(0, 0, 1)

Relevant centroid $R = (1, 0.5, 0)$

Non-relevant centroid $NR = (0.5, 0, 1.5)$

Quiz

1	✓	cat	(1, 0, 0)
2	✓	cat dog	(1, 1, 0)
3	✗	cat horse horse	(1, 0, 2)
4	✗	horse	(0, 0, 1)

Relevant centroid $R = (1, 0.5, 0)$

Non-relevant centroid $NR = (0.5, 0, 1.5)$

$Q = R + R - NR = (1.5, 1, -1.5)$

Quiz

1	✓	cat	(1, 0, 0)
2	✓	cat dog	(1, 1, 0)
3	✗	cat horse horse	(1, 0, 2)
4	✗	horse	(0, 0, 1)

$$Q = R + R - NR = (1.5, 1, -1.5)$$

$$\cos(Q, 1) = \frac{1.5 \cdot 1 + 1 \cdot 0 - 1.5 \cdot 0}{\sqrt{5.5} \cdot \sqrt{1}} = 0.64$$

$$\cos(Q, 3) = \frac{1.5 \cdot 1 + 1 \cdot 0 - 1.5 \cdot 2}{\sqrt{5.5} \cdot \sqrt{5}} = -0.29$$

Rocchio 1971 algorithm (SMART)

Used in practice:

$$\begin{aligned}\vec{q}_m &= \alpha \vec{q}_0 + \beta \mu(D_r) - \gamma \mu(D_{nr}) \\ &= \alpha \vec{q}_0 + \beta \frac{1}{|D_r|} \sum_{\vec{d}_j \in D_r} \vec{d}_j - \gamma \frac{1}{|D_{nr}|} \sum_{\vec{d}_j \in D_{nr}} \vec{d}_j\end{aligned}$$

- q_m : modified query vector; q_0 : original query vector; D_r and D_{nr} : sets of **known** relevant and nonrelevant documents respectively; α , β , and γ : weights
- New query moves towards relevant documents and away from nonrelevant documents.

Positive vs negative feedback

- Positive feedback is more valuable than negative feedback.
 - For example, set $\beta = 0.75$, $\gamma = 0.25$ to give higher weight to positive feedback.
- Many systems only allow positive feedback.
 - Why?
 - This is what we will do in assignment 3.

Relevance feedback - Assumption

- Relevance documents are similar
 - Term distribution in relevant documents will be similar
 - Term distribution in non-relevant documents will be different from those in relevant documents
 - Similarities between relevant and irrelevant documents are small

Violation of the assumption

- There are several clusters of relevant documents
- Examples:
 - Alternative terminology (Burma / Myanmar)
 - Disjunctive queries (*“Celebrities that use to work for Burger King”*)
 - Instances of general concepts(Feline → cat, tiger, lion, etc)

Relevance feedback: Problems

- Long queries are inefficient for typical IR engine.
 - Possible solution: Only reweight certain prominent terms
 - Perhaps top 20 by term frequency
- Users are often reluctant to provide explicit feedback
- It's often harder to understand why a particular document was retrieved after applying relevance feedback

Pseudo-relevance feedback

- Users are often reluctant to provide explicit feedback
- Pseudo-relevance feedback automates the “manual” part of true relevance feedback.
- Pseudo-relevance algorithm:
 - Retrieve a ranked list of hits for the user’s query
 - Assume that the top k documents are relevant.
 - Do relevance feedback (e.g., Rocchio)
- Works very well on average
 - But can go horribly wrong for some queries.
 - Several iterations can cause **query drift**.

Wildcard queries

Tolerant retrieval

- Spelling correction
 - “*see you on the wki*” → “*see you on the wiki*”
- Wildcard queries
 - “*colo*rful*” → “*colorful*”, “*colourful*”
- In both cases, the search engine needs to
 - construct the intended query (or queries)
 - compute the results for those queries (intersection, phrase, ranked retrieval)
 - list the results

Wildcard queries: one word

- **care***: find all docs containing any word beginning “care”.
- ***less**: find words ending in “less”
- **colo*r**: find all words beginning “colo” and ending in “r”
- general case: any numbers of ‘*’ placed anywhere in the word (we will **not** consider this case)
- special case: ‘*’ matches all words (you don’t need to consider this case)

Wildcard queries: several words

- **b* colo*r**: find all docs containing any word beginning “b”, **and** any word beginning with “colo” ending in “r”

Wildcard queries

- How do we find all words matching **care*** ?
- **Idea:** Go through all words in the vocabulary, and check which words match the regular expression **^care.***
 - e.g. using Java's regex library
- Would this work?

K-gram index

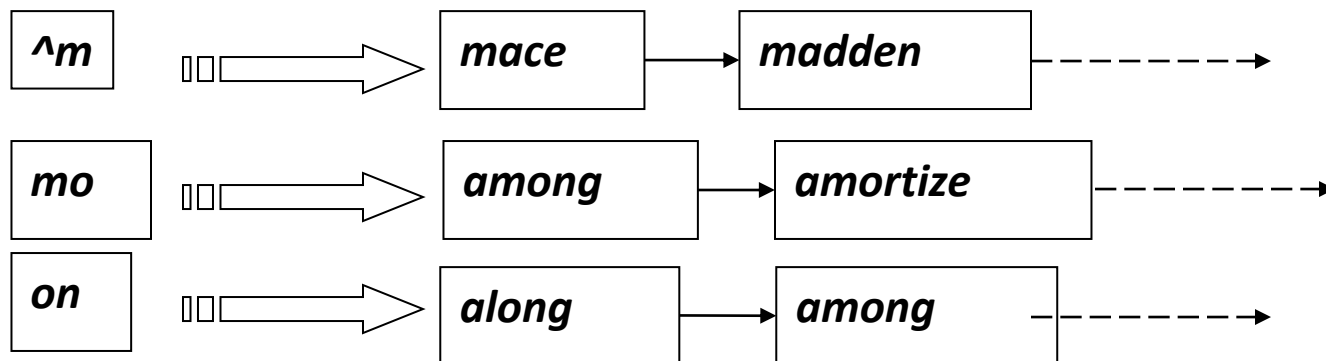
- For both wildcard queries and spelling correction we must **quickly** find words that
 - the user intended (for wildcard queries), or
 - the user probably intended (for spelling correction)
- A **k-gram index** is an index from k-grams (parts of words) to words.
 - **bigram** index when $k=2$
 - **trigram** index when $k=3$
 - etc.

K-grams

- The bigrams of **december**:
 - First add start and end symbol: **^december\$**
 - Bigrams are all two-letter sequences:
^d, de, ec, ce, em, mb, be, er, r\$
- The trigrams of **december**:
 - **^de, dec, ece, cem, emb, mbe, ber, er\$**
- A word of length n has $n+3-k$ k -grams


K-gram index

- For k-gram indexes, we can reuse a lot of ideas from our usual inverted indexes:
 - keys in a hashtable
 - values as arraylists



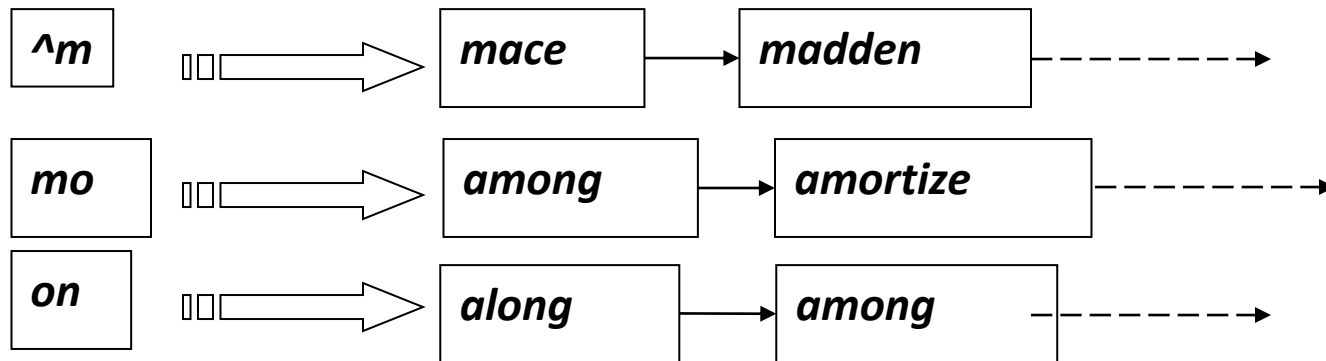
K-gram index and wildcards

- Suppose we want to find matches for **mon***
- How would you search?




K-gram index and wildcards

- Suppose we want to find matches for **mon***
- How would you search?
 - do an **intersection** search for ***^m mo on*** in k-gram index
 - **post-process the results using the regex library**
 - do a **union** search on the resulting words in the ordinary index



K-gram index and wildcards

- Suppose we want to find matches for **mon* tra***
- Which documents in this Venn diagram are you looking for?



Spelling correction

Spelling correction

- Two principal uses
 - Correcting document(s) being indexed
 - Correcting user queries to retrieve “right” answers
 - Usually documents are left intact, but queries spell-checked
- Two main flavors:
 - Isolated word
 - Will not catch typos resulting in correctly spelled words, e.g., *from* → *form*
 - Context-sensitive, e.g. *I flew form Heathrow to Narita.*

Spelling correction

- Fundamental premise – there is a lexicon from which the correct spellings come
- Two basic choices for this
 - (**Grammatical approach**) A standard lexicon such as
 - Webster's English Dictionary
 - An “industry-specific” lexicon – hand-maintained
 - (**Data-driven approach**) The lexicon of the indexed corpus
 - E.g., all words on the web
 - All names, acronyms etc.
 - (Including the mis-spellings)

Spelling correction of a single word

- What we will do in assignment 3:
 - Data-driven approach (no lexicon)
 - Assumption: A word is **not** misspelt if it appears in at least 1 document.
 - If a word has 0 occurrences it **might** be misspelt, and the search engine should suggest corrections

Spelling correction of a single word

- Methods:
 - **Edit distance**
 - **Weighted edit distance**
 - ***n*-gram overlap**
- These can (should?) be combined

Levenshtein (edit) distance

- What is $\text{dist}(\textit{intention}, \textit{execution})$?

i n t e n t i o n
n t e n t i o n
e t e n t i o n
e x e n t i o n
e x e n u t i o n
e x e c u t i o n

← delete *i*

← substitute *n* by *e*

← substitute *t* by *x*

← insert *u*

← substitute *n* by *c*

- Cost $1+2+2+1+2 = 8$
- Can be efficiently computed with dynamic programming

Computing Levenshtein distance

above broke

	#	a	b	o	v	e
#						
b						
r						
o						
k						
e						

Computing Levenshtein distance

above broke

	#	a	b	o	v	e
#	0	1	2	3	4	5
b	1					
r	2					
o	3					
k	4					
e	5					

Computing Levenshtein distance

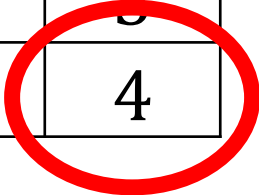
above broke

	#	a	b	o	v	e
#	0	1	2	3	4	5
b	1	2				
r	2					
o	3					
k	4					
e	5					

Computing Levenshtein distance

above broke

	#	a	b	o	v	e
#	0	1	2	3	4	5
b	1	2	1	2	3	4
r	2	3	2	3	4	5
o	3	4	3	2	3	4
k	4	5	4	3	4	5
e	5	6	5	4	5	4



Weighted edit distances

- As above, but the weight of an operation depends on the character(s) involved
 - Meant to capture OCR or keyboard errors, e.g. ***m*** more likely to be mis-typed as ***n*** than as ***q***
 - Therefore, replacing ***m*** by ***n*** is a smaller edit distance than by ***q***
 - This may be formulated as a probability model
- Requires weight matrix as input
- Modify dynamic programming to handle weights

Using edit distances

- Given a misspelt word in the query, find all words in the index within a preset edit distance (e.g. 2)
 1. Show terms you found to user as suggestions, ***or***
 2. Look up all possible corrections in our inverted index and return all docs ... slow, ***or***
 3. Run with a single most likely correction
- In assignment 3, we will opt for alternative 1.

Using edit distances

- Given a query, do we compute its edit distance to every dictionary term?
 - Expensive and slow
- How do we find the candidate dictionary terms?
 - One alternative: k-gram overlap
 - Use the k-gram index again!
 - Can also be used by itself for spelling correction

k -gram overlap

- Enumerate all the k -grams in the query string as well as in the lexicon

- *november*: ^no, nov, ove, vem, emb, mbe, ber, er\$

- *december*: ^de, dec, ece, cem, emb, mbe, ber, er\$

- overlap 4/12 unique trigrams

- the **Jacquard coefficient** = $4/12 = 0.33$

- generally, $\frac{|X \cap Y|}{|X \cup Y|}$ where X,Y are sets

Spelling correction of a single word

- E.g. **wki**
 - Do a union search in the k-gram index for
^w wk ki i\$
 - Calculate the Jacqard coefficient between **wki** and each of the resulting words
 - If the JC > some threshold for word *w*, calculate the Levenshtein distance between *w* and **wki**
 - If Levenshtein distance < some other threshold, then *w* is a potential correction
 - Add *w* to the list of corrections

Spelling correction, multi-word queries

- E.g. "**See yuoq on the wki**"
- Includes two (possibly) misspelled words: **yuoq** and **wki** with 0 postings
- Construct the lists of spelling suggestions for each word
 - Lists for words with > 0 postings will only contain themselves
 - Then merge the lists

Spelling correction, multi-word queries

- See youq on the wki

see	you	on	the	wiki
	your			ki
	youd			wi
	yous			wk
	youn			waki

- Now the lists have to merged to produce final suggestions

Spelling correction, multi-word queries

- Final list of suggestions (for instance):
 - See you on the wiki
 - See you on the ki
 - See you on the wi
 - See you on the wk
 - See you on the waki

General issues in spell correction

- We enumerate several possible alternatives to misspelled queries – which ones should we present to the user?
- Use heuristics:
 - The alternative matching most documents (expensive)
 - **The alternative likely to match most documents** (using heuristics, cheaper)
 - Query log analysis – what have others been searching for?
What has this user been searching for?

Query expansion

Query expansion

- improve retrieval results by adding synonyms / related terms to the query
- query-independent, "global" method

Why are synonyms important?

- As an example consider query q : [aircraft] . . .
 - . . . and document d containing “plane”, but not containing “aircraft”
 - A simple IR system will not return d for q .
 - Even if d is the most relevant document for q !
- We want to change this:
 - Return relevant documents even if there is no term match with the (original) query


Query expansion

- In global query expansion, the query is modified based on some global resource, i.e. a resource that is **not query-dependent**.
- Main information we use: (near-)**synonymy**
- A publication or database that collects (near)-S synonyms is called a **thesaurus**.
- We will look at two types of thesauri: manually created and automatically created.


Thesaurus-based query expansion

- For each term t in the query, expand the query with words the thesaurus lists as semantically related with t .
 - E.g. CARDIAC \rightarrow HEART
 - Generally increases recall
- May significantly decrease precision, particularly with ambiguous terms
 - INTEREST RATE \rightarrow INTEREST RATE HOBBY
- Widely used in specialized search engines for science and engineering
- It's very expensive to create a manual thesaurus and to maintain it over time.


PubMed: Manually curated thesaurus

 NCBI [Resources](#) [How To](#)

[Sign in to NCBI](#)

 **MeSH**

US National Library of Medicine
National Institutes of Health [Advanced](#) [Help](#)



PubMed

PubMed comprises more than 28 million citations for biomedical literature from MEDLINE, life science journals, and online books. Citations may include links to full-text content from PubMed Central and publisher web sites.

Using PubMed

- [PubMed Quick Start Guide](#)
- [Full Text Articles](#)
- [PubMed FAQs](#)
- [PubMed Tutorials](#)
- [New and Noteworthy](#)

PubMed Tools

- [PubMed Mobile](#)
- [Single Citation Matcher](#)
- [Batch Citation Matcher](#)
- [Clinical Queries](#)
- [Topic-Specific Queries](#)

More Resources

- [MeSH Database](#)
- [Journals in NCBI Databases](#)
- [Clinical Trials](#)
- [E-Utilities \(API\)](#)
- [LinkOut](#)


Automatic thesaurus construction

- Attempt to generate a thesaurus automatically by analyzing the distribution of words in documents
- Definition 1: Two words are similar if **they co-occur with similar words**.
 - “car” \approx “motorcycle” because both occur with “road”, “gas” and “license”, so they must be similar.
- Definition 2: Two words are similar if **they occur in a given grammatical relation** with the same words.
 - You can harvest, peel, eat, prepare, etc. apples and pears, so apples and pears must be similar.
- Co-occurrence is more robust, grammatical relations are more accurate.

Word embedding approaches

- Mapping words to vectors of real numbers
- If w_1 and w_2 have similar meaning, then $\text{vec}(w_1)$ and $\text{vec}(w_2)$ are similar
- Many approaches exist:
 - Latent Semantic Analysis (LSA), Random Indexing, Word2Vec (2013), Glove (2014), Fasttext (2017), Elmo (2018)...
 - Vectors have typically 50-300 dimensions
 - Words with similar semantics can be retrieved with a Nearest Neighbor software package

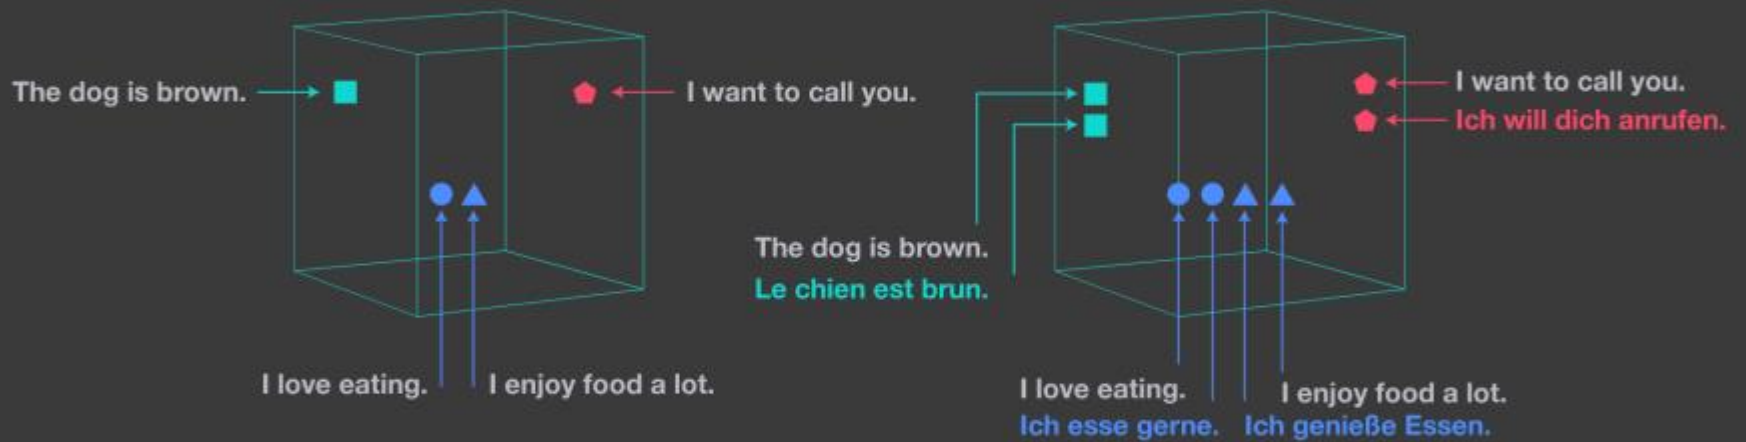
Word embedding approaches



Sentence embeddings

- Mapping sentences to vectors of real numbers
- If s_1 and s_2 have similar meaning, then $\text{vec}(s_1)$ and $\text{vec}(s_2)$ are similar
- Universal Sentence Encoder (2018), BERT (2018)
- LASER (2019) Multi-lingual sentence embeddings

LASER



Query expansion using query logs

- Example 1: After issuing the query [herbs], users frequently search for [herbal remedies].
 - → “herbal remedies” is potential expansion of “herb”.
- Example 2: Users searching for [flower pix] frequently click on the URL photobucket.com/flower. Users searching for [flower clipart] frequently click on the same URL.
 - → “flower clipart” and “flower pix” are potential expansions of each other.

Summary

- Ways of improving recall in search:
 - **Relevance feedback** (assignment 3.1)
 - **Wildcard queries** (3.3, 3.4)
 - **Spelling correction** (3.5, 3.6)
 - **Query expansion**